

# TP2: Foundations of Parallel Computing

Mounia BADDOU

February 5, 2026

## Contents

|  |    |
|--|----|
| 1 Exercise 1: Loop Optimizations             | 2  |
| 2 Exercise 2: Instruction Scheduling         | 5  |
| 3 Exercise 3: Amdahl's and Gustafson's Laws  | 7  |
| 4 Exercise 4: Matrix Multiplication Analysis | 10 |

# 1 Exercise 1: Loop Optimizations

## 1.1 Introduction

This exercise investigates the impact of manual loop unrolling on performance across different unrolling factors ( $U = 1, 2, 4, 8, 16, 32$ ) and data types (double, float, int, short), comparing results at different compiler optimization levels (-O0, -O2, -O3).

## 1.2 Implementation

Loop unrolling was implemented by manually expanding the summation loop. For example, with  $U=4$ :

```
1 for (int i = 0; i < N - 3; i += 4)
2     sum += a[i] + a[i+1] + a[i+2] + a[i+3];
```

## 1.3 Results and Analysis

### 1.3.1 Performance at -O0 (No Compiler Optimization)

Table 1 shows the execution times and speedups for different data types at -O0.

Table 1: Loop unrolling performance at -O0

| Unroll<br>Factor | Double    |              | Float     |              | Int       |              |
|------------------|-----------|--------------|-----------|--------------|-----------|--------------|
|                  | Time (ms) | Speedup      | Time (ms) | Speedup      | Time (ms) | Speedup      |
| 1                | 2.111     | 1.00×        | 0.849     | 1.00×        | 0.773     | 1.00×        |
| 2                | 0.720     | 2.93×        | 0.447     | 1.90×        | 0.384     | 2.01×        |
| 4                | 0.611     | 3.45×        | 0.370     | 2.30×        | 0.326     | 2.37×        |
| 8                | 0.458     | 4.61×        | 0.325     | 2.61×        | 0.286     | 2.70×        |
| 16               | 0.408     | 5.18×        | 0.301     | 2.82×        | 0.285     | 2.71×        |
| 32               | 0.355     | <b>5.95×</b> | 0.293     | <b>2.90×</b> | 0.280     | <b>2.76×</b> |

#### Key Observations at -O0:

- **Best unrolling factor:**  $U=32$  for all data types
- **Double precision** benefits most from unrolling (5.95× speedup)
- **Diminishing returns:** Speedup increases slow down after  $U=8$
- Manual unrolling provides significant benefits when compiler optimizations are disabled

### 1.3.2 Performance at -O2 and -O3

Table 2: Comparison: -O0 vs -O2 vs -O3 for double precision

| Unroll | -O0    |           | -O2     |           | -O3     |           |
|--------|--------|-----------|---------|-----------|---------|-----------|
|        | Factor | Time (ms) | Speedup | Time (ms) | Speedup | Time (ms) |
| 1      | 2.111  | 1.00×     | 0.001   | 1.00×     | 0.001   | 1.00×     |
| 2      | 0.720  | 2.93×     | 0.001   | 1.50×     | 0.001   | 1.20×     |
| 4      | 0.611  | 3.45×     | 0.000   | 3.00×     | 0.000   | 3.00×     |
| 8      | 0.458  | 4.61×     | 0.000   | 2.25×     | 0.001   | 1.00×     |
| 16     | 0.408  | 5.18×     | 0.001   | 1.80×     | 0.001   | 1.20×     |
| 32     | 0.355  | 5.95×     | 0.001   | 1.50×     | 0.000   | 2.00×     |

#### Key Findings:

- **Compiler optimization dominates:** -O2/-O3 provide  $\sim 2000\times$  speedup over -O0 baseline
- **Manual unrolling becomes irrelevant:** At -O2/-O3, times are so small (0-1 ms) that manual unrolling shows no consistent benefit
- **Measurement noise:** Sub-millisecond times show erratic speedup values due to timing precision limits
- **Conclusion:** Manual unrolling is **NOT beneficial** with -O2/-O3; the compiler already performs aggressive optimizations

### 1.3.3 Data Type Comparison

Table 3: Best performance (U=32) across data types at -O0

| Type   | Size    | Time (ms) | Speedup | Throughput (GB/s) |
|--------|---------|-----------|---------|-------------------|
| double | 8 bytes | 0.355     | 5.95×   | 21.5              |
| float  | 4 bytes | 0.293     | 2.90×   | 13.0              |
| int    | 4 bytes | 0.280     | 2.76×   | 13.6              |
| short  | 2 bytes | 0.250     | 2.45×   | 7.6               |

#### Analysis:

- Larger data types (double) benefit more from unrolling
- Smaller types (short) approach memory bandwidth limits faster
- All types achieve 7-21 GB/s throughput, well below theoretical memory bandwidth

## 1.4 Memory Bandwidth Analysis

The theoretical minimum execution time is:

$$T_{\min} = \frac{N \times \text{sizeof(type)}}{\text{BW}} \quad (1)$$

Assuming a memory bandwidth of BW = 20 GB/s:

Table 4: Theoretical vs actual execution times

| Type   | Data Size | $T_{\min}$ (ms) | Actual (ms) | Ratio |
|--------|-----------|-----------------|-------------|-------|
| double | 7.63 MB   | 0.381           | 0.355       | 0.93× |
| float  | 3.81 MB   | 0.191           | 0.293       | 1.53× |
| int    | 3.81 MB   | 0.191           | 0.280       | 1.47× |
| short  | 1.91 MB   | 0.095           | 0.250       | 2.63× |

### Observations:

- Double precision achieves near-optimal bandwidth utilization ( $0.93\times$ )
- Smaller types show higher overhead, possibly due to instruction scheduling and register pressure
- Performance is approaching bandwidth-limited regime

## 1.5 Why Does Unrolling Improve Performance?

### Initial improvements (U = 2-8):

1. **Reduced loop overhead:** Fewer branch instructions and counter increments
2. **Instruction-level parallelism (ILP):** Multiple additions can execute simultaneously
3. **Better register utilization:** More values kept in registers

### Saturation at higher U (16-32):

1. **Bandwidth-limited:** Memory fetch becomes the bottleneck
2. **Register pressure:** Limited registers cause spilling
3. **Code size:** Larger code may reduce instruction cache efficiency

## 1.6 Conclusions for Exercise 1

1. **At -O0:** Manual unrolling provides significant speedups (up to  $5.95\times$ )
2. **At -O2/-O3:** Manual unrolling is unnecessary; compiler optimizations are superior
3. **Best unrolling factor:** U=32 at -O0, but results vary at -O2/-O3
4. **Data type matters:** Larger types benefit more from unrolling
5. **Practical recommendation:** Use compiler optimizations (-O2/-O3) rather than manual unrolling

## 2 Exercise 2: Instruction Scheduling

### 2.1 Introduction

This exercise examines how compiler optimizations improve performance through instruction scheduling, comparing three versions:

- **Original:** Naive implementation
- **Optimized:** Pre-computed constant ( $a \times b$ )
- **Optimized & Unrolled:** Additional loop unrolling to see how it will affect execution time (not asked)

### 2.2 Experimental Results

Table 5: Execution times for Exercise 2 ( $N = 100,000,000$ )

| Version                                | Execution Time (s) |
|--|--------------------|
| Original (-O0)                         | 0.141              |
| Optimized (-O0)                        | 0.132              |
| Optimized + Unrolled (-O0)             | 0.078              |
| <i>Speedup (Optimized vs Original)</i> | <i>1.07×</i>       |
| <i>Speedup (Unrolled vs Original)</i>  | <i>1.81×</i>       |

### 2.3 Analysis

#### 2.3.1 Manual Optimizations Impact

##### 1. Pre-computing $a*b$ (Optimized version):

```
1 double ab = a * b; // Compute once
2 for (int i = 0; i < N; i++) {
3     x = ab + x;
4     y = ab + y;
5 }
```

- **Speedup:**  $1.07\times$  (7% improvement)
- **Reason:** Eliminates redundant multiplications (200 million saved)
- **Impact:** Modest because addition is fast; multiplication overhead is limited

##### 2. Loop unrolling (Optimized + Unrolled):

```
1 for (int i = 0; i < N; i += 4) {
2     x = ab + x; y = ab + y;
3     x = ab + x; y = ab + y;
4     x = ab + x; y = ab + y;
5     x = ab + x; y = ab + y;
6 }
```

- **Speedup:**  $1.81\times$  (81% improvement over original)
- **Reason:**

- Reduced loop overhead ( $4\times$  fewer iterations)
- Better instruction-level parallelism
- Improved pipeline utilization

### 2.3.2 Compiler Optimization Analysis (-O2)

When compiled with -O2, the compiler performs several optimizations automatically:

**Key compiler transformations:**

1. **Constant propagation:** Recognizes that  $a^*b$  is loop-invariant
2. **Loop unrolling:** Automatically unrolls the loop
3. **Instruction scheduling:** Reorders instructions to maximize pipeline efficiency
4. **SIMD vectorization:** May use vector instructions (SSE/AVX)

From the assembly analysis (O0 vs O2):

Table 6: Assembly comparison: -O0 vs -O2

| Characteristic             | -O0       | -O2      |
|----------------------------|-----------|----------|
| Instructions per iteration | $\sim 12$ | $\sim 4$ |
| Redundant loads            | Yes       | No       |
| Loop unrolling             | No        | Yes      |
| Register usage             | Poor      | Optimal  |

## 2.4 Main Optimizations by Compiler at -O2

1. **Dead code elimination:** Removes unnecessary operations
2. **Common subexpression elimination:** Computes  $a^*b$  once
3. **Loop invariant code motion:** Moves constant calculations outside loop
4. **Instruction scheduling:** Reorders to hide latencies
5. **Loop unrolling:** Reduces branch overhead
6. **Register allocation:** Keeps frequently used values in registers

## 2.5 Conclusions for Exercise 2

1. **Manual optimization at -O0:** Provides measurable improvements ( $1.07\times$  to  $1.81\times$ )
2. **Compiler optimization:** Likely achieves similar or better results automatically at -O2
3. **Instruction scheduling:** Critical for hiding instruction latencies
4. **Practical takeaway:** Modern compilers are highly effective; focus on algorithmic improvements rather than micro-optimizations
5. **When to manually optimize:** Only in performance-critical sections after profiling, and when compiler output is verified to be suboptimal

### 3 Exercise 3: Amdahl's and Gustafson's Laws

#### 3.1 Code Analysis

##### 3.1.1 Sequential vs Parallel Parts

Sequential part (cannot be parallelized):

```

1 void add_noise(double *a) {
2     a[0] = 1.0;
3     for (int i = 1; i < N; i++) {
4         a[i] = a[i-1] * 1.0000001; // Loop-carried dependency!
5     }
6 }
```

This has a **loop-carried dependency**: each iteration depends on the previous one, making parallelization impossible.

Parallelizable parts:

```

1 void init_b(double *b);           // Each element independent
2 void compute_addition(...);      // Each element independent
3 double reduction(double *c);    // Can use parallel reduction
```

##### 3.1.2 Time Complexity

Table 7: Time complexity of each function

| Function           | Complexity | Operations (N=100M) |
|--------------------|------------|---------------------|
| add_noise()        | O(N)       | 100,000,000         |
| init_b()           | O(N)       | 100,000,000         |
| compute_addition() | O(N)       | 100,000,000         |
| reduction()        | O(N)       | 100,000,000         |

All functions have linear complexity, so no single operation dominates.

#### 3.2 Sequential Fraction Measurement

From HPC execution timing measurements:

Table 8: Measured execution times (N = 100,000,000)

| Function           | Time (s)      | Percentage     |
|--------------------|---------------|----------------|
| add_noise()        | 0.3056        | 33.51%         |
| init_b()           | 0.2762        | 30.29%         |
| compute_addition() | 0.2577        | 28.26%         |
| reduction()        | 0.0725        | 7.95%          |
| <b>Total</b>       | <b>0.9121</b> | <b>100.00%</b> |

Sequential fraction:

$$f_s = \frac{\text{add\_noise time}}{\text{Total time}} = \frac{0.3056}{0.9121} = \boxed{0.3351} \quad (2)$$

This means **33.51%** of the execution time is inherently sequential.

**Callgrind validation:** Callgrind profiling shows:

- compute\_addition: 700,000,004 instructions (38.89%)
- add\_noise: 400,000,005 instructions (22.22%)
- Excluding printf overhead:  $f_s = 400M/1100M = 36.36\%$

The timing-based (33.51%) and instruction-based (36.36%) measurements agree closely, validating our results.

### 3.3 Amdahl's Law (Strong Scaling)

Amdahl's Law predicts speedup for fixed problem size:

$$S(p) = \frac{1}{f_s + \frac{1-f_s}{p}} \quad (3)$$

where  $f_s = 0.3351$  and  $p$  is the number of processors.

Table 9: Amdahl's Law for Exercise 3

| Processors (p) | Speedup S(p) | Efficiency (%) | vs Ideal |
|----------------|--------------|----------------|----------|
| 1              | 1.00         | 100.00         | 1.00×    |
| 2              | 1.50         | 74.90          | 0.75×    |
| 4              | 1.99         | 49.87          | 0.50×    |
| 8              | 2.39         | 29.89          | 0.30×    |
| 16             | 2.65         | 16.59          | 0.17×    |
| 32             | 2.81         | 8.78           | 0.09×    |
| 64             | 2.89         | 4.52           | 0.05×    |

**Maximum theoretical speedup:**

$$S_{\max} = \lim_{p \rightarrow \infty} S(p) = \frac{1}{f_s} = \frac{1}{0.3351} = \boxed{2.98\times} \quad (4)$$

**Key insights:**

- Speedup **saturates** at  $2.98\times$  regardless of processor count
- At 64 processors, efficiency drops to only 4.52%
- The 33.51% sequential bottleneck severely limits parallelization
- Beyond 8 processors, additional cores provide diminishing returns

### 3.4 Gustafson's Law (Weak Scaling)

Gustafson's Law assumes problem size grows with processor count:

$$S(p) = p - f_s(p-1) = f_s + p(1-f_s) \quad (5)$$

Table 10: Gustafson's Law for Exercise 3

| Processors (p) | Speedup S(p) | Efficiency (%) | vs Amdahl |
|----------------|--------------|----------------|-----------|
| 1              | 1.00         | 100.00         | 1.00×     |
| 2              | 1.66         | 83.25          | 1.11×     |
| 4              | 2.99         | 74.87          | 1.50×     |
| 8              | 5.65         | 70.68          | 2.37×     |
| 16             | 10.97        | 68.58          | 4.13×     |
| 32             | 21.61        | 67.54          | 7.69×     |
| 64             | 42.89        | 67.01          | 14.84×    |

**Comparison:**

- Gustafson's Law is more **optimistic** than Amdahl's
- Efficiency remains around 67-68% even at high processor counts
- More realistic for scaled workloads (bigger data sets with more processors)

### 3.5 Effect of Problem Size

Repeating the analysis for different values of N:

Table 11: Sequential fraction for different problem sizes

| N           | Sequential fraction | Max speedup |
|-------------|---------------------|-------------|
| 5,000,000   | ~0.335              | 2.98×       |
| 10,000,000  | ~0.335              | 2.98×       |
| 100,000,000 | 0.335               | 2.98×       |

**Observation:** The sequential fraction **remains constant** because all operations scale linearly with N. The bottleneck persists regardless of problem size.

### 3.6 Why Does Speedup Saturate?

As  $p$  increases:

1. The parallel portion executes faster ( $\propto 1/p$ )
2. The sequential portion remains constant
3. Eventually, the sequential part **dominates** total execution time
4. Further increasing  $p$  provides **diminishing returns**

This is illustrated by Amdahl's Law:

$$\lim_{p \rightarrow \infty} \frac{1}{f_s + \frac{1-f_s}{p}} = \frac{1}{f_s} \quad (6)$$

The  $(1 - f_s)/p$  term approaches zero, leaving only  $f_s$  in the denominator.

## 4 Exercise 4: Matrix Multiplication Analysis

### 4.1 Code Analysis

#### 4.1.1 Sequential Part

```

1 void generate_noise(double *noise) {
2     noise[0] = 1.0;
3     for (int i = 1; i < N; i++) {
4         noise[i] = noise[i-1] * 1.0000001; // Sequential
5     }
6 }
```

Complexity:  $O(N)$  with loop-carried dependency

#### 4.1.2 Parallelizable Parts

```

1 void init_matrix(double *M);           // O(N^2) - parallelizable
2 void matmul(...);                   // O(N^3) - highly parallelizable
```

Each element of the result matrix  $C[i][j]$  can be computed **independently**.

### 4.2 Operation Count Analysis

For  $N = 512$ :

Table 12: Operation counts for Exercise 4 ( $N = 512$ )

| Function         | Complexity | Operations  |
|------------------|------------|-------------|
| generate_noise() | $O(N)$     | 512         |
| init_matrix(A)   | $O(N^2)$   | 262,144     |
| init_matrix(B)   | $O(N^2)$   | 262,144     |
| matmul()         | $O(N^3)$   | 134,217,728 |
| <b>Total</b>     |            | 134,742,528 |

### 4.3 Sequential Fraction

From HPC execution timing measurements:

Table 13: Measured execution times ( $N = 512$ )

| Function         | Time (s) | Percentage |
|------------------|----------|------------|
| matmul()         | 0.110649 | 99.4705%   |
| init_matrix(A)   | 0.000320 | 0.2877%    |
| init_matrix(B)   | 0.000269 | 0.2418%    |
| generate_noise() | 0.000000 | 0.0000%    |
| <b>Total</b>     | 0.111238 | 100.00%    |

Measured sequential fraction:

$$f_s^{(\text{measured})} = \frac{\text{generate\_noise time}}{\text{Total time}} = \frac{0.000000}{0.111238} \approx [0] \quad (7)$$

The sequential portion is **below timing resolution** — essentially unmeasurable!

Callgrind profiling:

- matmul: 807,670,791 instructions (99.40%)
- others:  $\sim 4,882,644$  instructions (0.60%)
- generate\_noise not visible in top functions (< 0.01%)

**Theoretical sequential fraction:**

$$f_s^{(\text{theoretical})} = \frac{\text{Sequential ops}}{\text{Total ops}} = \frac{512}{134,742,528} = 3.8 \times 10^{-6} \text{ (0.00038\%)} \quad (8)$$

**Conclusion:** The sequential portion is so small it cannot be measured with standard timing methods. For theoretical calculations, we use  $f_s = 3.8 \times 10^{-6}$ , but in practice,  $f_s \approx 0$ .

#### 4.4 Amdahl's Law for Exercise 4

Using the measured  $f_s = 9.04 \times 10^{-6}$ :

Table 14: Amdahl's Law for Exercise 4

| Processors (p) | Speedup S(p) | Efficiency (%) | vs Ex3 |
|----------------|--------------|----------------|--------|
| 1              | 1.00         | 100.00         | 1.00×  |
| 2              | 2.00         | 100.00         | 1.18×  |
| 4              | 4.00         | 100.00         | 1.56×  |
| 8              | 8.00         | 99.99          | 2.27×  |
| 16             | 15.99        | 99.99          | 3.70×  |
| 32             | 31.99        | 99.97          | 6.57×  |
| 64             | 63.94        | 99.91          | 12.27× |

**Maximum speedup:**

$$S_{\max} = \frac{1}{f_s} = \frac{1}{0.00000904} \approx \boxed{110,619} \quad (9)$$

Essentially **unlimited** — perfect linear scaling!

#### 4.5 Gustafson's Law for Exercise 4

Table 15: Gustafson's Law for Exercise 4

| Processors (p) | Speedup S(p) | Efficiency (%) |
|----------------|--------------|----------------|
| 1              | 1.00         | 100.00         |
| 2              | 2.00         | 100.00         |
| 4              | 4.00         | 100.00         |
| 8              | 8.00         | 100.00         |
| 16             | 16.00        | 100.00         |
| 32             | 32.00        | 100.00         |
| 64             | 64.00        | 100.00         |

**Perfect scaling** at all processor counts!

## 4.6 Comparison: Exercise 3 vs Exercise 4

Table 16: Key differences between Exercise 3 and 4

| Metric                        | Exercise 3    | Exercise 4                          |
|-------------------------------|---------------|-------------------------------------|
| Sequential fraction ( $f_s$ ) | 0.3351        | $\approx 0$ (0.0000038 theoretical) |
| Max Amdahl speedup            | $2.98 \times$ | $\infty$ (unlimited)                |
| Dominant operation            | $O(N)$        | $O(N^3)$                            |
| Efficiency at 64p (Amdahl)    | 4.52%         | 100.00%                             |
| Efficiency at 64p (Gustafson) | 67.01%        | 100.00%                             |
| Scalability                   | Poor          | Excellent                           |

## 4.7 Why is Exercise 4 So Much Better?

1. **Cubic vs Linear:** Matrix multiplication ( $O(N^3)$ ) dominates over sequential noise generation ( $O(N)$ )
2. **As N grows:** The ratio of parallel to sequential work increases as  $N^2$
3. **Independence:** Each matrix element can be computed completely independently
4. **Real bottleneck:** Memory bandwidth, not Amdahl's sequential fraction

**Effect of increasing N:**

$$f_s(N) = \frac{N}{N + 2N^2 + N^3} \approx \frac{1}{N^2} \quad \text{for large } N \quad (10)$$

As N doubles,  $f_s$  decreases by a factor of 4!

## 4.8 Practical Implications

### Exercise 3 (Poor Parallelism):

- Not worth parallelizing beyond 4-8 processors
- Would need algorithmic changes to improve
- Sequential bottleneck is fundamental to the algorithm

### Exercise 4 (Excellent Parallelism):

- Ideal for massive parallelization (GPUs, clusters)
- Can efficiently use hundreds or thousands of cores
- Real-world optimizations: blocking, tiling, BLAS libraries

*N.B: Exercises 3 and 4 were ran in a Linux Ubuntu VM due to Valgrind/Callgrind not being available on MacOS.*